

Piyush Hinduja

Salt Lake City, UT | (385) 507-6190 | piyushinduja20@gmail.com | [Portfolio](#) | [LinkedIn](#)

EDUCATION

Master of Science in Computer Science

University of Utah

GPA: **3.91**/4.0

2023 – 2025

Bachelor of Engineering in Computer Engineering

University of Mumbai

GPA: **3.85**/4.0

2019 – 2023

SKILLS

Languages: *Expert (5+ yrs):* Python, Java, JavaScript *Intermediate (3+ yrs):* C, C++ *Beginner (1+ yrs):* Kotlin, SQL, Shell
Machine Learning: LLMs, PyTorch, TensorFlow, Scikit-Learn, BERT models, EasyOCR, HuggingFace, Langchain
Data Processing: RESTful APIs, Pandas, NumPy, PySpark, Databricks, OCR Pipelines, Matplotlib, Seaborn
Applications and Databases: Streamlit, Bash, Preswald, MongoDB, MySQL, Amazon RDS, Firebase, ChromaDB
Cloud & Miscellaneous: AWS (EC2, S3, RDS), GCP, Azure, Docker, Kubernetes, CloudLab, Git, CI/CD, Agile/SCRUM, JIRA

WORK EXPERIENCE

Data Analyst

University of Utah

March 2024 – Present

Salt Lake City, Utah

- Implemented advanced NLP analysis using **encoder-only model FinBERT, PyTorch, and Transformers** for financial document processing, **improving the prediction accuracy from 66% to 94.13%**
- Reduced per-filing **processing time by 75% using multi-cores and multi-thread environments** to parallelize file processing, API requests, and text cleaning operations
- Developed robust **prompt engineering** solutions to integrate LLMs & GenAIs (OpenAI, Anthropic, and Gemini) for automated financial text analysis, comparing AI-generated insights against traditional algorithms
- Optimized financial data processing workflows, achieving **faster execution and lower memory consumption** through optimized data structures and parallelized computations

Machine Learning Researcher

FLUX Research Group

August 2024 – April 2025

Salt Lake City, Utah

- Engineered automated deployment system using **Python and Linux Bash** to streamline **NVIDIA/CUDA toolkit** integration on high-performance computing clusters, reducing overhead by **90%**
- Analyzing wireless spectrum data to uncover scientific insights and **optimize power distribution across 20+ campus routing devices**

Full Stack Developer Intern

Exposys Data Labs

June 2021 – July 2021

Bangalore, India

- Spearheaded the launch of a tourism project, integrating key features to **enhance user experience** and deliver valuable information about 50+ Indian travel destinations
- Utilized **MERN stack** for an efficient frontend and backend integration, creating a responsive, user-friendly interface and enabling content management

PROJECTS

SEC Comment Letter Analysis | LLMs, OCR, Claude API

- Engineered a rolling-window algorithm to process and bundle SEC comment letter conversations, analyzing data from **50,000+ companies** via the EDGAR database
- Designed an **OCR pipeline using the EasyOCR library** to extract text from PDFs with **95% accuracy**, processing 2-page documents in under 2 seconds
- Constructed prompt-based labeling system using **Claude API** to classify letter content with **98.25% accuracy**, integrating metadata for contextual inference

Road Damage Detection | Object-Detection, YOLOv5, FasterRCNN

[GitHub Link](#)

- Collaborated with a team of four to develop a road damage detection system using **YOLOv5 and FasterRCNN**, integrating a user-friendly web interface (deployed via **Streamlit**), leveraging the RDD-2020 dataset
- Achieved model accuracies of **52.67% with YOLOv5** and **44.45% with FasterRCNN** by training and fine-tuning on different hyper-parameters
- The system plays a vital role, delivering a **70% improvement in road damage detection accuracy** over manual methods, helping lower the frequency of accidents and fatalities

Dependency Parser | PyTorch, GloVe Embeddings

[GitHub Link](#)

- Architected an end-to-end **Transition-Based Dependency Parser in PyTorch**, implementing a multi-class classifier for predicting parsing actions based on stack and buffer states
- Designed modular pipeline incorporating data preprocessing, model architecture, training workflow, and evaluation metrics
- Integrated **GloVe word embeddings** for enhanced semantic representation, contributing to **UAS of 0.76 and LAS of 0.705**

Language Modeling | NLP, RNN, N-gram, PyTorch

[GitHub Link](#)

- Developed **character-level LSTM model** in PyTorch to predict character sequences based on **50-character context windows**
- Implemented teacher forcing during training to prevent error propagation, using ground truth characters instead of predictions for subsequent timesteps
- Achieved **perplexity of 9.43** through systematic model architecture optimization and training methodology